

# //////////parallel dfs and bfs

```
#include <bits/stdc++.h>
# include <omp.h>
using namespace std;

// Function to perform Parallel BFS
void parallelBFS(vector<vector<int>>& adj_list, int source, vector<bool>& visited, vector<int>& bfs_order) {
    queue<int> q;
    q.push(source);

    // Parallel loop over the queue
    #pragma omp parallel
    {
        while (!q.empty()) {
            // Get the next vertex from the queue
            #pragma omp for
            for (int i = 0; i < q.size(); i++) {
                int curr = q.front();
                q.pop();

                // If the current vertex has not been visited, mark it as visited
                // and explore all its neighbors
                if (!visited[curr]) {
                    #pragma omp critical
                    {
                        visited[curr] = true;
                        bfs_order.push_back(curr); // add the visited node to the bfs_order vector
                    }
                    for (int j = 0; j < adj_list[curr].size(); j++) {
                        int neighbor = adj_list[curr][j];

                        // Add the neighbor to the queue if it has not been visited
                        if (!visited[neighbor]) {
                            q.push(neighbor);
                        }
                    }
                }
            }
        }
    }
}
```

```

// Function to perform Parallel DFS
void parallelDFS(vector<vector<int>>& adj_list, int source, vector<bool>& visited, vector<int>&
dfs_order) {
    stack<int> s;
    s.push(source);

    // Parallel loop over the stack
    #pragma omp parallel
    {
        while (!s.empty()) {
            // Get the next vertex from the stack
            #pragma omp for
            for (int i = 0; i < s.size(); i++) {
                int curr = s.top();
                s.pop();

                // If the current vertex has not been visited, mark it as visited
                // and explore all its neighbors
                if (!visited[curr]) {
                    #pragma omp critical
                    {
                        visited[curr] = true;
                        dfs_order.push_back(curr); // add the visited node to the dfs_order vector
                    }
                    for (int j = 0; j < adj_list[curr].size(); j++) {
                        int neighbor = adj_list[curr][j];

                        // Add the neighbor to the stack if it has not been visited
                        if (!visited[neighbor]) {
                            s.push(neighbor);
                        }
                    }
                }
            }
        }
    }

    int main() {
        // Construct the adjacency list
        vector<vector<int>> adj_list = {
            {1, 2},

```

```
{0, 3, 4},  
{0, 5, 6},  
{1},  
{1},  
{2},  
{2}  
};
```

```
// Perform Parallel BFS from node 0  
int source = 0;  
int n = adj_list.size();  
vector<bool> visited(n, false);  
vector<int> bfs_order;  
parallelBFS(adj_list, source, visited, bfs_order);
```

```
// Print the visited nodes and the BFS order  
cout << "BFS order: ";  
for (int i = 0; i < bfs_order.size(); i++) {  
    cout << bfs_order[i] << " ";  
}  
cout << endl;
```

```
for (int i = 0; i < n; i++) {  
    if (visited[i]) {  
        cout << "Node " << i << " has been visited" << endl;  
    }  
}
```

```
// Perform Parallel DFS from  
// reset the visited vector  
fill(visited.begin(), visited.end(), false);  
vector<int> dfs_order;  
parallelDFS(adj_list, source, visited, dfs_order);
```

```
// Print the visited nodes and the DFS order  
cout << "DFS order: ";  
for (int i = 0; i < dfs_order.size(); i++) {  
    cout << dfs_order[i] << " ";  
}  
cout << endl;
```

```
for (int i = 0; i < n; i++) {  
    if (visited[i]) {
```

```

        cout << "Node " << i << " has been visited" << endl;
    }
}

return 0;
}

```

```

////////////////////////////////////

```

### **Parallel mergesort**

```

#include <iostream>

```

```

#include <cstdlib>

```

```

#include <ctime>

```

```

#include <omp.h>

```

```

using namespace std;

```

```

void merge(int arr[], int l, int m, int r) {

```

```

    int n1 = m - l + 1;

```

```

    int n2 = r - m;

```

```

    int L[n1], R[n2];

```

```

    for (int i = 0; i < n1; i++) {

```

```

        L[i] = arr[l + i];

```

```

    }

```

```

    for (int j = 0; j < n2; j++) {

```

```

        R[j] = arr[m + 1 + j];

```

```

    }

```

```

    int i = 0, j = 0, k = l;

```

```

    while (i < n1 && j < n2) {

```

```

        if (L[i] <= R[j]) {

```

```

            arr[k] = L[i];

```

```

            i++;

```

```

        } else {

```

```

            arr[k] = R[j];

```

```

            j++;

```

```

        }

```

```

        k++;

```

```

    }

```

```

while (i < n1) {
    arr[k] = L[i];
    i++;
    k++;
}

while (j < n2) {
    arr[k] = R[j];
    j++;
    k++;
}
}

void mergeSort(int arr[], int l, int r) {
    if (l < r) {
        int m = l + (r - l) / 2;
        #pragma omp parallel sections
        {
            #pragma omp section
            mergeSort(arr, l, m);
            #pragma omp section
            mergeSort(arr, m + 1, r);
        }
        merge(arr, l, m, r);
    }
}

int main() {
    srand(time(nullptr));
    const int size = 10000;
    int arr[size];

    for (int i = 0; i < size; i++) {
        arr[i] = rand() % 10000;
    }

    double start = omp_get_wtime();
    mergeSort(arr, 0, size - 1);
    double end = omp_get_wtime();

    cout << "Sorted array: " << endl;
    for (int i = 0; i < size; i++) {
        cout << arr[i] << " ";
    }
}

```

```

    }
    cout << endl;
    cout << "Time taken: " << end - start << " seconds" << endl;

    return 0;
}

```

////////////////////////////////////

### **Parallel bubble sort**

```
#include <stdio.h>
```

```
#include <omp.h>
```

```

void bubble_sort(int arr[], int n) {
    int i, j;
    for (i = 0; i < n - 1; i++) {
        if (i % 2 == 0) {
            #pragma omp parallel for shared(arr)
            for (j = 0; j < n - 1; j += 2) {
                if (arr[j] > arr[j+1]) {
                    int temp = arr[j];
                    arr[j] = arr[j+1];
                    arr[j+1] = temp;
                }
            }
        }
        else {
            #pragma omp parallel for shared(arr)
            for (j = 1; j < n - 1; j += 2) {
                if (arr[j] > arr[j+1]) {
                    int temp = arr[j];
                    arr[j] = arr[j+1];
                    arr[j+1] = temp;
                }
            }
        }
    }
}

```

```

int main() {
    int arr[] = {64, 34, 25, 12, 22, 11, 90};
    int n = sizeof(arr) / sizeof(arr[0]);

    bubble_sort(arr, n);
}

```

```

printf("Sorted array: ");
for (int i = 0; i < n; i++)
    printf("%d ", arr[i]);
printf("\n");

return 0;
}
//////////min, max , average and sum in parallel reduction

```

```

#include <iostream>

#include <vector>

#include <omp.h>

using namespace std;

void min_reduction(vector<int>& arr) {
    int min_value = INT_MAX;

    #pragma omp parallel for reduction(min: min_value)

    for (int i = 0; i < arr.size(); i++) {
        if (arr[i] < min_value) {
            min_value = arr[i];
        }
    }

    cout << "Minimum value: " << min_value << endl;
}

```

```

void max_reduction(vector<int>& arr) {

```

```

int max_value = INT_MIN;

#pragma omp parallel for reduction(max: max_value)

for (int i = 0; i < arr.size(); i++) {

    if (arr[i] > max_value) {

        max_value = arr[i];

    }

}

cout << "Maximum value: " << max_value << endl;

}

```

```

void sum_reduction(vector<int>& arr) {

    int sum = 0;

    #pragma omp parallel for reduction(+: sum)

    for (int i = 0; i < arr.size(); i++) {

        sum += arr[i];

    }

    cout << "Sum: " << sum << endl;

}

```

```

void average_reduction(vector<int>& arr) {

    int sum = 0;

    #pragma omp parallel for reduction(+: sum)

    for (int i = 0; i < arr.size(); i++) {

        sum += arr[i];

    }

}

```



```

    }

    cout << "Average: " << (double)sum / arr.size() << endl;
}

```

```

int main() {

    vector<int> arr = {5, 2, 9, 1, 7, 6, 8, 3, 4};

    min_reduction(arr);

    max_reduction(arr);

    sum_reduction(arr);

    average_reduction(arr);

}

```

#### **////////// cuda addition of two large vectors**

```

#include <iostream>
#include <cuda_runtime.h>
#include <bits/stdc++.h>
// Kernel function for vector addition
__global__ void vectorAdd(const float* a, const float* b, float* c, int size)
{
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx < size)
        c[idx] = a[idx] + b[idx];
}

```

```

int main()
{
    int size = 1000000; // Size of the vectors
    size_t bytes = size * sizeof(float);

    // Allocate memory on the host (CPU)
    float* h_a = new float[size];
    float* h_b = new float[size];
    float* h_c = new float[size];
}

```

```

// Initialize input vectors
for (int i = 0; i < size; ++i) {
    h_a[i] = i;
    h_b[i] = i;
}

// Allocate memory on the device (GPU)
float* d_a, * d_b, * d_c;
cudaMalloc((void**)&d_a, bytes);
cudaMalloc((void**)&d_b, bytes);
cudaMalloc((void**)&d_c, bytes);

// Copy input data from host to device
cudaMemcpy(d_a, h_a, bytes, cudaMemcpyHostToDevice);
cudaMemcpy(d_b, h_b, bytes, cudaMemcpyHostToDevice);

// Define block and grid sizes
int threadsPerBlock = 256;
int blocksPerGrid = (size + threadsPerBlock - 1) / threadsPerBlock;

// Launch kernel on the GPU
vectorAdd<<<blocksPerGrid, threadsPerBlock>>>(d_a, d_b, d_c, size);

// Copy result from device to host
cudaMemcpy(h_c, d_c, bytes, cudaMemcpyDeviceToHost);

// Print the first 10 elements of the result
for (int i = 0; i < 10; ++i) {
    std::cout << h_c[i] << " ";
}
std::cout << std::endl;

// Free memory
delete[] h_a;
delete[] h_b;
delete[] h_c;
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);

return 0;
}

```

# **////////// cuda matrix multiplication**

```
#include <iostream>
```

```
#include <cstdlib>
```

```
#include <bits/stdc++.h>
```

```
// CUDA kernel for matrix multiplication
```

```
__global__ void matrixMultiply(int *a, int *b, int *c, int N)
```

```
{  
    int row = blockIdx.y * blockDim.y + threadIdx.y;  
    int col = blockIdx.x * blockDim.x + threadIdx.x;
```

```
    if (row < N && col < N) {  
        int sum = 0;  
        for (int k = 0; k < N; ++k) {  
            sum += a[row * N + k] * b[k * N + col];  
        }  
        c[row * N + col] = sum;  
    }
```

```
}
```

```
int main()
```

```
{
```

```
    int N = 4; // Matrix size
```

```
    int *a, *b, *c; // Host matrices
```

```
    int *d_a, *d_b, *d_c; // Device matrices
```

```
    int matrixSize = N * N * sizeof(int);
```

```
    // Allocate host memory
```

```
    a = (int*)malloc(matrixSize);
```

```
    b = (int*)malloc(matrixSize);
```

```
    c = (int*)malloc(matrixSize);
```

```
    // Initialize host matrices
```

```
    for (int i = 0; i < N * N; ++i) {
```

```
        a[i] = i + 1;
```

```
        b[i] = i + 1;
```

```
    }
```

```
    // Allocate device memory
```

```

cudaMalloc((void**)&d_a, matrixSize);
cudaMalloc((void**)&d_b, matrixSize);
cudaMalloc((void**)&d_c, matrixSize);

// Transfer data from host to device
cudaMemcpy(d_a, a, matrixSize, cudaMemcpyHostToDevice);
cudaMemcpy(d_b, b, matrixSize, cudaMemcpyHostToDevice);

// Define block and grid dimensions
dim3 threadsPerBlock(2, 2);
dim3 blocksPerGrid((N + threadsPerBlock.x - 1) / threadsPerBlock.x,
                  (N + threadsPerBlock.y - 1) / threadsPerBlock.y);

// Launch kernel
matrixMultiply<<<blocksPerGrid, threadsPerBlock>>>(d_a, d_b, d_c, N);

// Transfer results from device to host
cudaMemcpy(c, d_c, matrixSize, cudaMemcpyDeviceToHost);

// Print result
for (int i = 0; i < N * N; ++i) {
    std::cout << c[i] << " ";
    if ((i + 1) % N == 0)
        std::cout << std::endl;
}

// Free memory
free(a);
free(b);
free(c);
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);

return 0;
}

```

**Or**

```

#define N 16
#include <bits/stdc++.h>
#include <cuda_runtime.h>

```

```

Using namespace std;
__global__ void matrixMult (int *a, int *b, int *c, int width);
int main() { int a[N][N], b[N][N], c[N][N];
    int *dev_a, *dev_b, *dev_c;
    // initialize matrices a and b with appropriate values
    int size = N * N * sizeof(int);
    cudaMalloc((void **) &dev_a, size);
    cudaMalloc((void **) &dev_b, size);
    cudaMalloc((void **) &dev_c, size);
    cudaMemcpy(dev_a, a, size, cudaMemcpyHostToDevice);
    cudaMemcpy(dev_b, b, size, cudaMemcpyHostToDevice);
    dim3 dimGrid(1, 1);
    dim3 dimBlock(N, N);
    matrixMult<<dimGrid, dimBlock>>(dev_a, dev_b, dev_c, N);
    cudaMemcpy(c, dev_c, size, cudaMemcpyDeviceToHost);
    for (int i = 0; i < N * N; ++i) {
        std::cout << c[i] << " ";
        if ((i + 1) % N == 0)
            std::cout << std::endl;
    }
    cudaFree(dev_a);
    cudaFree(dev_b);
    cudaFree(dev_c);
__global__ void matrixMult (int *a, int *b, int *c, int width)
{ int k, sum = 0;
    int col = threadIdx.x + blockDim.x * blockIdx.x;
    int row = threadIdx.y + blockDim.y * blockIdx.y;
    if(col < width && row < width)
    { for (k = 0; k < width; k++)
        sum += a[row * width + k] * b[k * width + col];
        c[row * width + col] = sum;
    }
}
}

```

<http://users.wfu.edu/choss/CUDA/docs/Lecture%205.pdf>

**nvcc program.cu -o program**

**Huffman**

**#include <iostream>**

**#include <cuda\_runtime.h>**

**\_\_global\_\_ void buildHuffmanTree(int\* frequencies, int\* tree, int n) {**

```

int i = threadIdx.x + blockIdx.x * blockDim.x;
if (i < n) {
    // Find the two lowest frequency nodes
    int min1 = INT_MAX, min2 = INT_MAX;
    int minIndex1, minIndex2;
    for (int j = 0; j < n; j++) {
        if (frequencies[j] != 0 && frequencies[j] < min1) {
            min2 = min1;
            minIndex2 = minIndex1;
            min1 = frequencies[j];
            minIndex1 = j;
        } else if (frequencies[j] != 0 && frequencies[j] < min2) {
            min2 = frequencies[j];
            minIndex2 = j;
        }
    }
    // Combine the two lowest frequency nodes into a new node
    int newNodeIndex = n + i;
    frequencies[newNodeIndex] = min1 + min2;
    tree[newNodeIndex] = 0;
    tree[newNodeIndex + n] = 0;
    if (minIndex1 < minIndex2) {
        tree[newNodeIndex] = minIndex1;
        tree[newNodeIndex + n] = minIndex2;
    } else {
        tree[newNodeIndex] = minIndex2;
        tree[newNodeIndex + n] = minIndex1;
    }
}
}

```

```

int main() {
    int n = 256;
    int* frequencies;
    int* tree;
    cudaMalloc(&frequencies, n * sizeof(int));
    cudaMalloc(&tree, 2 * n * sizeof(int));

    // Initialize frequencies
    for (int i = 0; i < n; i++) {
        frequencies[i] = i + 1;
    }
}

```

```

int numBlocks = (n + 255) / 256;
buildHuffmanTree<<<numBlocks, 256>>>(frequencies, tree, n);

// Encode the data using the Huffman tree
// ...

cudaFree(frequencies);
cudaFree(tree);
return 0;
}

```

Huffman encoding;

```

#include <iostream>
#include <queue>
#include <vector>

```

// Node structure for the Huffman tree

```

struct Node {
    char data;
    unsigned frequency;
    Node* left;
    Node* right;

    Node(char data, unsigned frequency)
        : data(data), frequency(frequency), left(nullptr), right(nullptr) {}

    ~Node() {
        delete left;
        delete right;
    }
};

```

// Comparison function for priority queue

```

struct Compare {
    bool operator()(Node* left, Node* right) {
        return left->frequency > right->frequency;
    }
};

```

// Kernel function for generating Huffman codes on the GPU

```

__global__ void generateCodesKernel(Node* root, char* codes, int* codeLengths, int

```

```

codesSize) {
    int tid = threadIdx.x + blockIdx.x * blockDim.x;

    if (tid < codesSize) {
        Node* node = root;
        int codeIndex = tid * codesSize;
        int codeLength = 0;

        while (node) {
            if (node->left && tid < node->left->frequency) {
                node = node->left;
                codes[codeIndex + codeLength] = '0';
            } else if (node->right) {
                node = node->right;
                codes[codeIndex + codeLength] = '1';
            } else {
                break;
            }
            codeLength++;
        }

        codeLengths[tid] = codeLength;
    }
}

```

**// Huffman encoding function**

```

void huffmanEncodeGPU(const char* input, char* output, int size, const char* codes,
const int* codeLengths, int codesSize) {
    char* d_input;
    char* d_output;
    char* d_codes;
    int* d_codeLengths;

```

**// Allocate device memory**

```

cudaMalloc((void**)&d_input, size * sizeof(char));
cudaMalloc((void**)&d_output, size * codesSize * sizeof(char));
cudaMalloc((void**)&d_codes, codesSize * codesSize * sizeof(char));
cudaMalloc((void**)&d_codeLengths, codesSize * sizeof(int));

```

**// Copy input data to device memory**

```

cudaMemcpy(d_input, input, size * sizeof(char), cudaMemcpyHostToDevice);

```

**// Copy Huffman codes to device memory**



```

    cudaMemcpy(d_codes, codes, codesSize * codesSize * sizeof(char),
cudaMemcpyHostToDevice);

    // Copy code lengths to device memory
    cudaMemcpy(d_codeLengths, codeLengths, codesSize * sizeof(int),
cudaMemcpyHostToDevice);

    // Configure kernel execution parameters
    int blockSize = 256;
    int gridSize = (codesSize + blockSize - 1) / blockSize;

    // Launch the kernel to generate codes on the GPU
    generateCodesKernel<<<gridSize, blockSize>>>(root, d_codes, d_codeLengths,
codesSize);

    // Copy the encoded data from device to host memory
    cudaMemcpy(output, d_output, size * codesSize * sizeof(char),
cudaMemcpyDeviceToHost);

    // Free device memory
    cudaFree(d_input);
    cudaFree(d_output);
    cudaFree(d_codes);
    cudaFree(d_codeLengths);
}

int main() {
    std::string text = "Hello, world!";
    int size = text.size();

    // Count frequencies of characters
    std::vector<unsigned> frequencies(256, 0);
    for (char c : text) {
        frequencies[c]++;
    }

    // Create a priority queue to store nodes
    std::priority_queue<Node*, std::vector<Node*>, Compare> pq

```